

```

import gradio as gr
import torch
from transformers import AutoTokenizer, AutoModelForCausalLM
import PyPDF2
import io

# Load model and tokenizer
model_name = "ibm-granite/granite-3.2-2b-instruct"
tokenizer = AutoTokenizer.from_pretrained(model_name)
model = AutoModelForCausalLM.from_pretrained(
    model_name,
    torch_dtype=torch.float16 if torch.cuda.is_available() else torch.float32,
    device_map="auto" if torch.cuda.is_available() else None
)

if tokenizer.pad_token is None:
    tokenizer.pad_token = tokenizer.eos_token

def generate_response(prompt, max_length=1024):
    inputs = tokenizer(prompt, return_tensors="pt", truncation=True, max_length=
    max_length)

    if torch.cuda.is_available():
        inputs = {k: v.to(model.device) for k, v in inputs.items()}

    with torch.no_grad():
        outputs = model.generate(
            **inputs,
            max_length=max_length,
            temperature=0.7,
            do_sample=True,
            pad_token_id=tokenizer.eos_token_id
        )

    response = tokenizer.decode(outputs[0], skip_special_tokens=True)
    response = response.replace(prompt, "").strip()
    return response

def extract_text_from_pdf(pdf_file):
    if pdf_file is None:
        return ""

    try:
        pdf_reader = PyPDF2.PdfReader(pdf_file)
        text = ""
        for page in pdf_reader.pages:
            text += page.extract_text() + "\n"
        return text
    except Exception as e:
        return f"Error reading PDF: {str(e)}"

def requirement_analysis(pdf_file, prompt_text):
    # Get text from PDF or prompt
    if pdf_file is not None:
        content = extract_text_from_pdf(pdf_file)
        analysis_prompt = f"Analyze the following document and extract key soft
    else:
        analysis_prompt = f"Analyze the following requirements and organize the

    return generate_response(analysis_prompt, max_length=1200)

def code_generation(prompt, language):
    code_prompt = f"Generate {language} code for the following requirement:\n\n
    return generate_response(code_prompt, max_length=1200)

```

```

# Create Gradio interface
with gr.Blocks() as app:
    gr.Markdown("# AI Code Analysis & Generator")

    with gr.Tabs():
        with gr.TabItem("Code Analysis"):
            with gr.Row():
                with gr.Column():
                    pdf_upload = gr.File(label="Upload PDF", file_types=[".pdf"])
                    prompt_input = gr.Textbox(
                        label="Or write requirements here",
                        placeholder="Describe your software requirements...",
                        lines=5
                    )
                analyze_btn = gr.Button("Analyze")

            with gr.Column():
                analysis_output = gr.Textbox(label="Requirements Analysis")

            analyze_btn.click(requirement_analysis, inputs=[pdf_upload, prompt_input], outputs=[analysis_output])

        with gr.TabItem("Code Generation"):
            with gr.Row():
                with gr.Column():
                    code_prompt = gr.Textbox(
                        label="Code Requirements",
                        placeholder="Describe what code you want to generate..",
                        lines=5
                    )
                language_dropdown = gr.Dropdown(
                    choices=["Python", "JavaScript", "Java", "C++", "C#", "TypeScript", "Go", "Ruby", "Swift", "Kotlin"],
                    label="Programming Language",
                    value="Python"
                )
            generate_btn = gr.Button("Generate Code")

            with gr.Column():
                code_output = gr.Textbox(label="Generated Code", lines=20)

            generate_btn.click(code_generation, inputs=[code_prompt, language_dropdown], outputs=[code_output])

app.launch(share=True)

```



```
/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:9
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settir
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to acc
warnings.warn(
```

```
tokenizer_config.json:      8.88k/? [00:00<00:00, 345kB/s]
```

```
vocab.json:      777k/? [00:00<00:00, 7.76MB/s]
```

```
merges.txt:      442k/? [00:00<00:00, 13.7MB/s]
```

```
tokenizer.json:      3.48M/? [00:00<00:00, 42.9MB/s]
```

```
added_tokens.json: 100%                               87.0/87.0 [00:00<00:00, 4.24kB/s]
```

```
special_tokens_map.json: 100%                         701/701 [00:00<00:00, 26.9kB/s]
```

```
100% 100% 100% 100% 100% 100% 100% 100% 100% 100%
```